

Recession Forecasting - Yield Curve and ML

Mehmet Zahid Duran, Selim Kir, Alexander Schied

December 26, 2025

Abstract

This paper evaluates a recession forecasting pipeline designed to predict U.S. National Bureau of Economic Research (NBER) recessions twelve months ahead. The framework integrates information from the Treasury yield curve, macroeconomic indicators, labor market variables, credit spreads, and explicit term premium decompositions. Principal component analysis (PCA) is used to extract interpretable yield curve factors, whose stability is examined through rolling-window diagnostics. Forecasting models are evaluated under strict temporal discipline using expanding-window cross-validation and out-of-sample threshold selection. Results show that yield curve slope and curvature factors are highly informative for financially driven recessions such as the Global Financial Crisis, while macroeconomic variables dominate performance during abrupt shock episodes such as the COVID-19 recession. The findings highlight the regime-dependent nature of recession signals and the importance of separating expectations-driven yield movements from term premium distortions.

1 Introduction

Forecasting economic recessions is a central problem in macroeconomics and financial economics, particularly when predictions must be made at horizons that are relevant for policy and risk management. The Treasury yield curve has long been recognized as one of the most reliable leading indicators of U.S. recessions. However, the interpretation of yield curve signals has become more complex in recent decades due to the expansion of unconventional monetary policy, large-scale asset purchases, and sustained compression of term premia.

This paper develops and evaluates a forecasting pipeline that predicts U.S. recessions twelve months in advance. Unlike contemporaneous classification approaches, this design imposes a genuine forecasting challenge. The analysis integrates yield curve information across maturities, principal component representations, macroeconomic indicators, labor market measures, and explicit decompositions that separate expectations-driven yield movements from term-premium-driven dynamics. Special attention is paid to avoiding information leakage and to diagnosing the stability of yield curve factors across monetary regimes.

2 Data and Target Construction

The target variable is the monthly NBER recession indicator. To ensure a true forecasting horizon, the target is shifted backward by twelve months, such that all predictors available at time t are used to forecast recession status at time $t + 12$.

The yield curve dataset consists of monthly Treasury yields spanning maturities from three months to thirty years. These yields are used both directly and as inputs to principal component analysis. Macroeconomic indicators include the unemployment rate, year-over-year CPI inflation, year-over-year industrial production growth, and consumer sentiment. Labor market conditions are proxied using year-over-year growth in nonfarm payroll employment. Financial conditions are captured through the federal funds rate, corporate credit spreads, and an estimated ten-year Treasury term premium.

To distinguish expectations-driven yield curve movements from term premium effects,

the ten-year yield is decomposed into an expectations proxy and a term premium component. An adjusted slope measure is constructed by subtracting the term premium from the long rate prior to computing the slope relative to the two-year yield. This allows direct testing of whether recession signals arise from anticipated policy easing or from non-expectational forces such as quantitative easing and safe-asset demand.

3 Principal Component Analysis of the Yield Curve

3.1 Explained Variance

Figure 1 presents the scree plot of the PCA applied to standardized yields across maturities. The first principal component explains approximately 97% of total variance, while the second and third components explain roughly 2.5% and 0.1%, respectively. This indicates that the U.S. yield curve is overwhelmingly driven by a common level factor, although slope and curvature components remain economically meaningful despite their small variance shares.




Figure 1: Explained variance ratio of yield curve principal components.

Table 2 reports the explained variance ratios numerically.

Table 1: Explained variance ratio of principal components.

index	0
PC1	0.971
PC2	0.027
PC3	0.002

3.2 Factor Loadings

Figure 2 displays the PCA loading heatmap across maturities. The first principal component loads positively and nearly uniformly across all maturities, consistent with a level factor. The second component exhibits negative loadings at the short end and positive loadings at the long end, representing yield curve slope. The third component captures curvature, reflecting differential movements of the belly of the curve relative to short and long maturities.




Figure 2: Heatmap of yield curve PCA loadings by maturity.

Detailed loadings are reported in Table 4.

Table 2: PCA loadings for yield curve components.

index	PC1	PC2	PC3
y_3m	0.329	-0.446	0.514
y_6m	0.331	-0.401	0.218
y_1y	0.334	-0.314	-0.089
y_2y	0.337	-0.140	-0.366
y_3y	0.338	-0.026	-0.421
y_5y	0.337	0.158	-0.294
y_7y	0.335	0.267	-0.132
y_{10y}	0.333	0.370	0.098
y_{30y}	0.325	0.537	0.505

3.3 Factor Dynamics and Recessions

Figure 3 plots the time series of the first three principal components with NBER recession shading. The level factor reflects long-run secular declines in interest rates and major monetary regime shifts. The slope factor exhibits pronounced swings around tightening cycles and yield curve inversions preceding recessions. The curvature factor shows episodic spikes during periods of financial stress and policy intervention.




Figure 3: Time series of yield curve principal components with recession shading.

3.4 Rolling Loading Stability

Figures 4–6 present rolling-window PCA loading estimates for the first three components. While the overall factor structure is preserved, substantial drift is observed, particularly at the long end of the curve during the Global Financial Crisis and the COVID-19 period. These findings suggest that yield curve factors are regime-dependent and that long-maturity yields become structurally distorted during periods of unconventional monetary policy.




Figure 4: Rolling PCA loading drift for the level factor (PC1).




Figure 5: Rolling PCA loading drift for the slope factor (PC2).




Figure 6: Rolling PCA loading drift for the curvature factor (PC3).

4 Forecasting Models and Evaluation

Models are estimated using multiple feature sets, including yield spreads, raw yields, yield PCA factors, macroeconomic variables, and hybrid combinations incorporating credit spreads and term premium proxies. Classifiers include regularized logistic regression, random forests, gradient boosting machines, and extreme gradient boosting. Thresholds are selected using out-of-fold predicted probabilities generated exclusively from the training sample via expanding-window cross-validation.

Table 6 reports holdout performance across multiple train–test splits.

5 Scenario Analysis: COVID-19 and the Global Financial Crisis

To assess regime dependence, models are evaluated separately on the COVID-19 recession (2019–2021) and the Global Financial Crisis (2007–2009). Figures 7 and 8 report F2 scores, emphasizing recall, while Figures 9 and 10 report precision–recall AUC.




Figure 7: Top F2 scores during the COVID-19 recession.

Table 3: Holdout sensitivity results across feature sets and models.

TestFrac	FeatureSet	Model	ThrObjective	ThrOpt	ThrObjectiveValue _{onTrainOOF}
0.200	spreads	Logistic	balanced _a cc	0.590	0.575
0.200	spreads	RandomForest	balanced _a cc	0.655	0.667
0.200	spreads	GradBoost	balanced _a cc	0.999	0.485
0.200	spreads	XGBoost	balanced _a cc	0.096	0.624
0.200	yield _p ca	Logistic	balanced _a cc	0.864	0.659
0.200	yield _p ca	RandomForest	balanced _a cc	0.101	0.638
0.200	yield _p ca	GradBoost	balanced _a cc	0.889	0.589
0.200	yield _p ca	XGBoost	balanced _a cc	0.086	0.598
0.300	spreads	Logistic	balanced _a cc	0.161	0.678
0.300	spreads	RandomForest	balanced _a cc	0.096	0.695
0.300	spreads	GradBoost	balanced _a cc	0.001	0.563
0.300	spreads	XGBoost	balanced _a cc	0.096	0.658
0.300	yield _p ca	Logistic	balanced _a cc	0.136	0.700
0.300	yield _p ca	RandomForest	balanced _a cc	0.096	0.668
0.300	yield _p ca	GradBoost	balanced _a cc	0.006	0.544
0.300	yield _p ca	XGBoost	balanced _a cc	0.146	0.673
0.400	spreads	Logistic	balanced _a cc	0.226	0.740
0.400	spreads	RandomForest	balanced _a cc	0.086	0.757
0.400	spreads	GradBoost	balanced _a cc	0.270	0.552
0.400	spreads	XGBoost	balanced _a cc	0.086	0.745
0.400	yield _p ca	Logistic	balanced _a cc	0.400	0.750
0.400	yield _p ca	RandomForest	balanced _a cc	0.141	0.768
0.400	yield _p ca	GradBoost	balanced _a cc	0.001	0.665
0.400	yield _p ca	XGBoost	balanced _a cc	0.275	0.727




Figure 8: Top F2 scores during the Global Financial Crisis.




Figure 9: Top PR-AUC scores during the COVID-19 recession.




Figure 10: Top PR-AUC scores during the Global Financial Crisis.

Table 8 summarizes scenario-specific results numerically.

6 Discussion

Results demonstrate that yield curve PCA factors, particularly slope and curvature, provide strong leading information for financially driven recessions such as the Global Financial Crisis. In contrast, macroeconomic variables dominate performance during sudden shock episodes like COVID-19, where traditional yield curve signals lag the downturn. Rolling PCA diagnostics reveal that factor stability deteriorates during periods of unconventional monetary policy, underscoring the importance of term premium controls and regime-aware modeling.

7 Conclusion

This paper shows that recession forecasting performance is highly regime-dependent. Yield curve factors remain powerful predictors, but their interpretation must account for term premium distortions and structural instability. Combining PCA-based yield representations with macroeconomic indicators and disciplined threshold selection yields robust and interpretable forecasts across diverse economic environments.

Table 4: Scenario-based evaluation results (GFC vs COVID).

Scenario	TrainEnd	TestStart	TestEnd	FeatureSet	Model	ThrObject
GFC _{test} _{2007_2009}	2006-12-01	2007-01-01	2009-12-01	spreads	Logistic	balanced,
GFC _{test} _{2007_2009}	2006-12-01	2007-01-01	2009-12-01	spreads	RandomForest	balanced,
GFC _{test} _{2007_2009}	2006-12-01	2007-01-01	2009-12-01	spreads	GradBoost	balanced,
GFC _{test} _{2007_2009}	2006-12-01	2007-01-01	2009-12-01	spreads	XGBoost	balanced,
GFC _{test} _{2007_2009}	2006-12-01	2007-01-01	2009-12-01	yield _{pc}	Logistic	balanced,
GFC _{test} _{2007_2009}	2006-12-01	2007-01-01	2009-12-01	yield _{pc}	RandomForest	balanced,
GFC _{test} _{2007_2009}	2006-12-01	2007-01-01	2009-12-01	yield _{pc}	GradBoost	balanced,
GFC _{test} _{2007_2009}	2006-12-01	2007-01-01	2009-12-01	yield _{pc}	XGBoost	balanced,
COVID _{test} _{2019_2021}	2018-12-01	2019-01-01	2021-12-01	spreads	Logistic	balanced,
COVID _{test} _{2019_2021}	2018-12-01	2019-01-01	2021-12-01	spreads	RandomForest	balanced,
COVID _{test} _{2019_2021}	2018-12-01	2019-01-01	2021-12-01	spreads	GradBoost	balanced,
COVID _{test} _{2019_2021}	2018-12-01	2019-01-01	2021-12-01	spreads	XGBoost	balanced,
COVID _{test} _{2019_2021}	2018-12-01	2019-01-01	2021-12-01	yield _{pc}	Logistic	balanced,
COVID _{test} _{2019_2021}	2018-12-01	2019-01-01	2021-12-01	yield _{pc}	RandomForest	balanced,
COVID _{test} _{2019_2021}	2018-12-01	2019-01-01	2021-12-01	yield _{pc}	GradBoost	balanced,
COVID _{test} _{2019_2021}	2018-12-01	2019-01-01	2021-12-01	yield _{pc}	XGBoost	balanced,